

data cleaning principles

Karl Broman

Biostatistics & Medical Informatics, UW–Madison

@kwbroman

kbroman.org

github.com/kbroman

kbroman.org/Talk_DataCleaning



These are slides for a talk for the csv,conf,v6 (<https://csvconf.com/>) on May 4–5, 2021.

Data analysts spend a lot of time organizing and cleaning data, but few of us have been trained to do so. Why is that?

Some say that data cleaning is difficult to generalize. But I think there are some general principles. Moreover, I think we have an important shared experience in data cleaning that we can commiserate about, and through which we can learn from each other.

Tidy data are all alike,
but every messy dataset
is messy in its own way.

– Hadley Wickham

2

Hadley's talking more about data organization than data cleanliness. And his point is that if you make data tidy, it simplifies all the downstream analyses.

But **is** every messy dataset **uniquely** messy?

For sure, many my collaborators have shown impressive creativity in their approach to organizing and managing data. But we do see many of the same sorts of problems over and over.

If I clean up [Medicare] data ...
does any of the knowledge I gain ...
apply to the processing of RNA-seq data?

– Roger Peng

3

In his discussion of David Donoho's paper about data science, Roger Peng wrote about how data cleaning is frustratingly difficult to generalize.

But my answer to his question is **absolutely!**

A person with experience cleaning one dataset has important experience to draw upon when moving to another dataset, even if it's of a totally different nature.

Data Mishaps Night

Join us for the first inaugural Data Mishaps Night!
We will feature a lineup of data mistake stories with
a focus on the human aspect of data work and
lessons learned the hard way.



Caitlin Hudon & Laura Ellis
dataMishapsNight.com

4

In February, 2021, Caitlin Hudon and Laura Ellis organized a Friday evening conference where 16 people gave short presentations on data mishaps.

Many of the stories concerned mistakes in data cleaning, and these seemed to bring out a strong sense of shared experience. We have suffered and struggled through very similar data problems.

Data cleaning

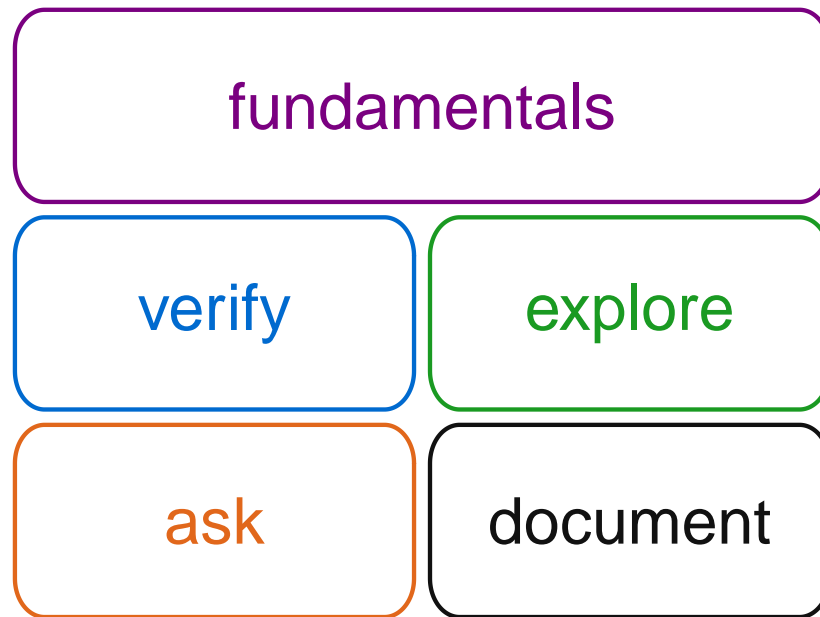
- ▶ tedious
- ▶ embarrassing
- ▶ needs context
- ▶ doesn't feel like progress
- ▶ requires creativity
- ▶ requires coding prowess
- ▶ source of most problems

5

Really, I think we don't usually teach data cleaning because it's something we prefer to keep private.

We're shy about it.

And data cleaning code is our ugliest code.



I'm proposing a set of basic principles for data cleaning, and splitting them into five groups. There are some fundamental principles, followed by four basic ideas: verify things that you expect, explore to find further oddities, ask questions, and document what you've done.

fundamentals

1. Don't clean data when you're tired or hungry.

(paraphrasing Ghazal Gulati)

7

At her talk at the Data Mishaps night, Ghazal Gulati emphasized this point, of not cleaning data when you're tired or hungry.

Data cleaning requires considerable concentration, and you need to allow sufficient time to do the work. If you're in a hurry, you'll miss things.

fundamentals

2. Don't trust anyone (even yourself)

“my motto is ‘trust no one’
...except maybe @kwbroman?”

– Jenny Bryan

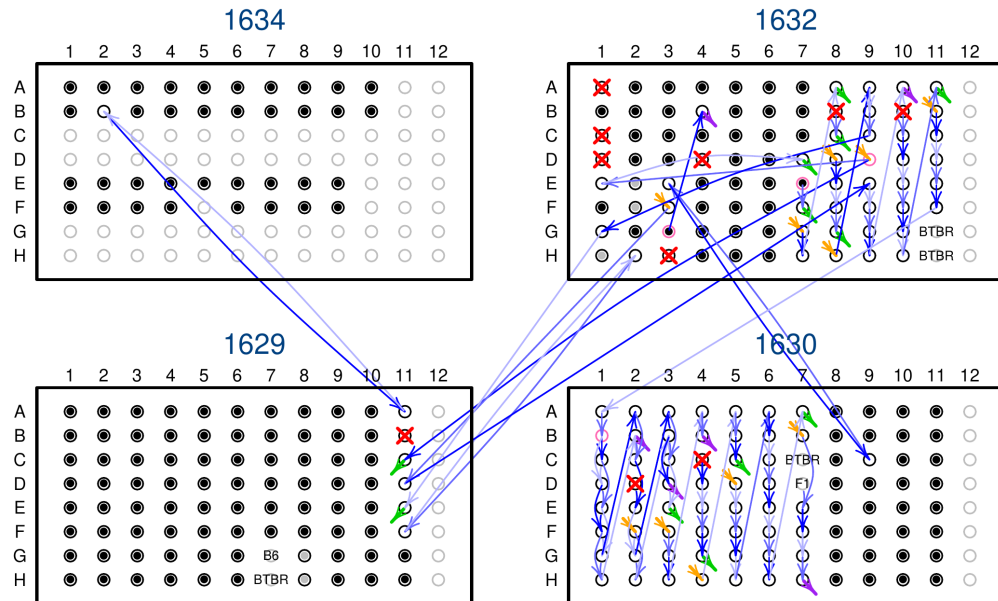
8

Next: don't trust anyone. Even if the initial data cleaning was done by someone you respect, you should double-check things that they may have missed. And data cleaning is an ongoing process.

Jenny Bryan's tweet is among the nicest things anyone has said about me.

fundamentals

3. Think about what might have gone wrong and how it might be revealed



9

Personally, I think this is the most important principle for data cleaning. It has been central in guiding my approach.

The figure here is an illustration of the most startlingly result I've had in data cleaning: a genetics project where 25% of the DNA samples had been mislabelled. The samples were arranged in wells in 8×12 plates; four of the six plates are shown here. The dots indicate the correct DNA was placed in the well, but the arrows point from where a sample should have been to where it actually was placed.

I ultimately came to this finding by thinking about what might have gone wrong in the project, checking for particularly problems, and then following the trail of evidence to this mess.

fundamentals

4. Use care in merging

	A	B	C	D	E	F	G		
1	id	glucose.0	glucose.5	glucose.15	glucose.30	insulin.0	insulin.5		
2	DO-221	145.742786	206.452638	216.640608	299.55501	0.74455	2.0264		
3	DO-222								
4	DO-223								
5	DO-224								
6	DO-225								
7	DO-226								
8	DO-227								
9	DO-228								
10	DO-229								
11	DO-230								

10

Many problems arise due to mistakes when merging data from multiple files. A common problem is a change in the data arrangement, such as in the order of columns.

Focus on the labels (which are more likely correct), rather than the position of variables in a file (which are more likely to change).

5. Dates & categories suck

11

The fifth fundamental principle is that dates and categories suck. You'll expend an inordinate amount of time dealing with these: typos in category labels, different date formats, people who died a decade they were born or lived to be 150.

Just be glad if you're not dealing with time zones.

But you may be asking yourself, "How is this a principle."

Principle:

a fundamental truth that guides our thinking

12

Yeah, are these principles? I was thinking the same thing. Was I drifting away from principles and more to just stuff to know or do?

This seems a pretty good definition, and is sufficiently broad to cover what I'm proposing, for the most part.

fundamentals

5. Dates & categories suck

13

So yeah, this counts as a principle.

Much of your pain will come from the dates and categorical data; you should be ready for that.

verify

6. Check that distinct things are distinct

	A	B	C	D	E	F	G
1	WiscID	ID	NEOID	Fem_CA	Fem_lmax	Fem_lmin	Fem_J
2	F2.C1W.F.1248	1248	NEO183	0.7524	0.1427	0.1006	0.2433
3	F2.C1W.M.1250	1250	NEO184	0.7669	0.1556	0.09652	0.2521
4	F2.C1W.F.1251	1251	NEO185	0.7613	0.1549	0.09659	0.2515
5	F2.C1W.F.1254	1254	NEO186	0.7475	0.1503	0.08603	0.2363
6	F2.C1W.M.1257	1257	NEO187	0.8197	0.1849	0.1056	0.2905
7	F2.___.F.715	715	NEO764	0.6017	0.09662	0.05969	0.1563
8	F2.___.F.751	751	NEO765	0.7273	0.1304	0.08735	0.2178
9	F2.___.F.1251	1251	NEO766	0.6675	0.1157	0.07814	0.1938
10	F2.___.M.1340	1340	NEO768	0.6656	0.1387	0.08122	0.2199
11	F2.C1W.M.739	739	NEO779	0.9336	0.2828	0.1628	0.4456

14

The next major section of principles concern efforts to **verify** all of the things that should be true for your data.

First up: are the variables that are supposed to have distinct values really showing distinct values? Here, there are a pair of individual identifiers that are duplicated.

verify

7. Check that matching things match

	A	B	C	D
1	id	sex	n_gen	age_days
2	F20.25	M	20	75
3	F21.30	M	21	75
4	F21.68	M	21	71
5	F22.52	M	22	73
6	F21.71	F	22	63
7	F22.116	F	22	57
8	F21.F20.9.M5	M	20	82
9	F21.F20.18.M5	M	20	77
10	F20.26	M	20	75
11	F21.62	M	21	72



	A	B	C	D
1	id	sex	age_at_dosing	n_gen
2	F22.69	F	67	22
3	F22.106	F	69	22
4	F22.70	F	67	22
5	F22.107	F	69	22
6	F21.71	F	65	21
7	F22.116	F	62	22
8	F22.73	F	65	22
9	F22.117	F	62	22
10	F21.108	F	62	21
11	F22.118	F	59	22

15

Next, are the things that are supposed to match actually matching? This often concerns data that are repeated in multiple data files. For example, you might have body weight included multiple times; are individuals' weights the same in all files? Here, there's a column for "number of generations" that shows an inconsistency between two files.

verify

8. Check calculations



16

Any time there is a calculation, you should verify the values. This is useful both for ensuring that you understand the calculation, and you also have the chance to reduce round-off error (though that seldom matters). Occasionally you may see mistakes.

Here, I'm plotting HOMA-IR which is the product of blood glucose and blood insulin; the values in the data on the x-axis and the values I calculated on the y-axis. It's often better to look at the differences directly; in the right panel I plot the differences on the y-axis. It appears that there is a batch of individuals that whose values were rounded more coarsely.

verify


9. Look for other instances of a problem

17

As with software testing, any time you find a problem, be sure to look for other instances of that problem.

explore

10. Make lots of plots



18

A particularly important aspect of data cleaning concerns just exploring the data to identify potential problems. And a particularly important aspect of that exploration is to just make lots of plots.


Plot variables over time or by subject ID, which may indicate things like batch effects.

Make scatterplots of variables against each other, looking particularly for outliers. Outliers could be real biological variation, but they could also be data entry problems, like a pair of numbers being swapped, or the weights entered in grams rather than milligrams.

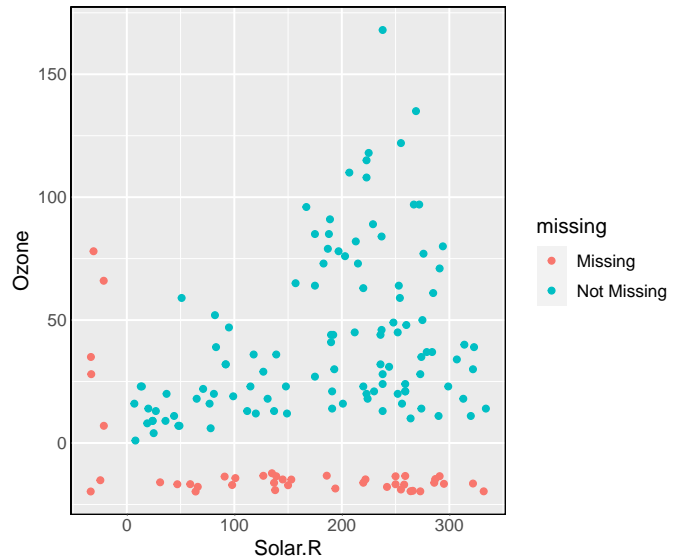
explore

11. Look at missing value patterns

{visdat}



{naniar}



19


The pattern of missing data can be particularly informative about problems. Two particularly useful tools are the R packages {visdat} and {naniar}. The latter provides scatterplots that include the cases that are missing one or both variables.

visdat: <https://docs.ropensci.org/visdat/>

naniar: <http://naniar.njtierney.com/>

explore

12. With massive data, make more plots not fewer



20

With large-scale datasets, it can be hard to make the sort of exploratory plots that you'd typically make. With oodles of data, you'd think you'd be looking at oodles of plots, but there's a tendency to give up and not look at any.


It's hard to look at 500 histograms, but it can be done. Superimpose a bunch of density estimates, maybe highlighting some portion of them. You can also pull out a couple of summary statistics, such as the median and inter-quartile range.

Or here I'm looking at the equivalent of 500 boxplots. I sorted a set of gene expression microarrays by their median, and then plotted the median in blue, the 25th and 75th percentile in black, the 10th and 90th in green, 5th and 95th in red, and 1st and 99th in purple.

With these data, it became apparent that there were 120 badly behaved arrays, with median shifted to the right and with a long left tail.

explore

13. Follow up all artifacts



kbroman.org/blog/2012/04/25/microarrays-suck

21

Wow the clash of those colors is particularly bad.

This is a heat map of the correlation matrix for a set of gene expression microarrays. The plaid pattern was a shock to me, and was caused by a set of bad arrays that we hadn't noticed previously.

My point here is simply to follow up all artifacts.

If you see something weird, follow through and try to figure out the underlying cause. It could be an error, or a set of bad assays, or it could be the most interesting finding in the study.

ask

- 14. Ask questions
- 15. Ask for the primary data
- 16. Ask for metadata
- 17. Ask why data are missing

22

I'm sure I'll be running out of time at this point, and so I've condensed these last two sections.

First, don't be shy about asking questions: about the data, to get the primary data, to also get metadata that explains the data (particularly variable names and such).

Among the most important questions to ask are why data are missing (for example, that an assay didn't work, that the value was too low, or that the value was too high).

document

- 18. Create checklists & pipelines
- 19. Document not just what but why
- 20. Expect to recheck

23

The final section of principles concerns documentation.

Create checklists and pipelines for yourself and others, so that when you return to similar data, you will remember many of the things to check, and you can build on what you've learned.

And for this sort of work, we need things to be **more** than reproducible; you will need to capture not just what you did, but also **why**. For example, if you decide to omit some subset of samples, will you remember 2 years from now just **why** you chose to omit them?

Finally, data cleaning is not a single step in the analysis chain; rather, it is an ongoing process that you will need to continually revisit as you delve deeper into the data. Keep an eye out for hints of problems, and arrange your work with the expectation that you'll need to re-run everything at some point.

fundamentals

1. Don't clean data when tired or hungry
2. Don't trust anyone (even yourself)
3. Think about what might have gone wrong
4. Use care in merging
5. Dates & categories suck

verify

6. Verify that distinct things are distinct
7. Verify that matching things match
8. Check calculations
9. Look for other instances of problems

explore

10. Make lots of plots
11. Look at missing value patterns
12. With big data make more plots
13. Follow up all artifacts

ask

14. Ask questions
15. Ask for the primary data
16. Ask for metadata
17. Ask why data are missing

document

18. Create checklists & pipelines
19. Document not just what but why
20. Expect to recheck

In summary, data cleaning is not an activity that needs to be constructed from scratch in each instance for each dataset. There are a number of principles that can guide our approach to cleaning data.

There may actually be more commonalities in our data cleaning experiences and methods than in the following stages of our data analysis work.

Slides: kbroman.org/Talk_DataCleaning



`kbroman.org`

`github.com/kbroman`

`@kwbroman`

Here's where you can find me and these slides.